


WASTE CLASSIFICATION

Building a model using pretrained model



ADEL YERKEBULAN
BDA-2203

TABLE OF CONTENTS

04

INTRODUCTION

08

CONCLUSION

05

DATASET

06


BASELINE MODEL

07

ENHANCED MODEL

INTRODUCTION

This project focuses on implementing a deep learning-based waste classification system using ResNet. The model aims to assist in automated sorting for waste management, classifying images into 30 distinct categories. The enhanced ResNet model achieved an accuracy of 84% on the test set.



DATASET

The dataset was sourced from the `waste_dataset/images/images` directory. It contains 30 classes, including aerosol cans, aluminum food cans, and plastic bags, among others. The data was split into:

- Training set: 70% of the data.
- Validation set: 15% of the data.
- Test set: 15% of the data.

Dataset Preprocessing

- Image Size: Resized to 224 × 224 pixels.

```
[ ] import os
import torchvision
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, random_split

data_dir = "../waste_dataset/images/images"

IMG_SIZE = (224, 224)
BATCH_SIZE = 32

transform = transforms.Compose([
    transforms.Resize(IMG_SIZE),
    transforms.ToTensor(),
    transforms.Normalize([0.5, 0.5, 0.5], [0.5, 0.5, 0.5])
])

dataset = datasets.ImageFolder(root=data_dir, transform=transform)

class_names = dataset.classes
print(f"Class names: {class_names}")

train_size = int(0.7 * len(dataset))
val_size = int(0.15 * len(dataset))
test_size = len(dataset) - train_size - val_size


train_dataset, val_dataset, test_dataset = random_split(dataset, [train_size, val_size, test_size])

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)
```

➡ Class names: ['aerosol_cans', 'aluminum_food_cans', 'aluminum_soda_cans', 'cardboard_boxes', 'cardboard_packagin

BASELINE MODEL

The base model was ResNet18, pre-trained on ImageNet. Its final fully connected layer was modified to output 30 classes



```
[ ] import torch
    import torch.nn as nn
    from torchvision import models

    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model = models.resnet18(pretrained=True)

    num_classes = len(dataset.classes)
    model.fc = nn.Linear(model.fc.in_features, num_classes)
    model = model.to(device)

    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```


```
➡ /usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight or path
  warnings.warn(msg)
```

MODEL TRAINING

- Loss Function: CrossEntropyLoss for multi-class classification.
- Optimizer: Adam with learning rates:
- Base model: 0.001.
- Batch Size: 32.
- Epochs: 5.

```
▶ train_model(model, train_loader, val_loader, loss, optimizer, 5)
```

↔	Average loss: 1.571679, Train accuracy: 0.540381, Val accuracy: 0.557778
	Average loss: 0.991221, Train accuracy: 0.693429, Val accuracy: 0.633333
	Average loss: 0.698026, Train accuracy: 0.776381, Val accuracy: 0.700889
	Average loss: 0.532077, Train accuracy: 0.828476, Val accuracy: 0.712000
	Average loss: 0.455369, Train accuracy: 0.840286, Val accuracy: 0.708000




ENHANCED MODEL

The enhanced model includes additional regularization via a dropout layer and a fully connected intermediate layer with ReLU activation

```
[ ] class EnhancedResNet(nn.Module):
    def __init__(self, pretrained_model, num_classes):
        super(EnhancedResNet, self).__init__()
        self.features = nn.Sequential(*list(pretrained_model.children())[:-1])
        self.fc = nn.Sequential(
            nn.Flatten(),
            nn.Linear(pretrained_model.fc.in_features, 512),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(512, num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        x = self.fc(x)
        return x

enhanced_model = EnhancedResNet(model, num_classes).to(device)
optimizer = torch.optim.Adam(enhanced_model.parameters(), lr=0.0001)
```



MODEL TRAINING

- Loss Function: CrossEntropyLoss for multi-class classification.
- Optimizer: Adam with learning rates:
- Enhanced model: 0.0001.
- Batch Size: 32.
- Epochs: 5.
- Learning Rate Scheduler: Decayed learning rate during training to improve convergence.

```
▶ train_model(enhanced_model, train_loader, val_loader, loss, optimizer, 5)
```

⇒	Average loss: 1.055853, Train accuracy: 0.786000, Val accuracy: 0.807111
	Average loss: 0.189688, Train accuracy: 0.943429, Val accuracy: 0.819556
	Average loss: 0.123070, Train accuracy: 0.957619, Val accuracy: 0.814667
	Average loss: 0.100828, Train accuracy: 0.963238, Val accuracy: 0.816889
	Average loss: 0.087163, Train accuracy: 0.967619, Val accuracy: 0.817333

RESULTS

Baseline model

Enhanced mode

Both models were evaluated on the test set. Detailed metrics are provided below

Baseline model:

- Precision (Macro Average): 75%
- Recall (Macro Average): 72%
- F1-Score (Macro Average): 72%
- Accuracy: 72%.


Classification Metrics

- Precision (Macro Average): 84%
- Recall (Macro Average): 83%
- F1-Score (Macro Average): 83%
- Accuracy: 84%.

	precision	recall	f1-score	support
0	0.92	0.74	0.82	62
1	0.36	0.37	0.36	73
2	0.73	0.70	0.72	63
3	0.62	0.37	0.46	63
4	0.48	0.66	0.56	74
5	0.52	0.79	0.63	63
6	0.85	0.85	0.85	65
7	1.00	0.64	0.78	78
8	0.75	0.99	0.85	81
9	0.72	0.84	0.78	83
10	0.62	0.88	0.73	78
11	0.94	0.73	0.82	83
12	0.77	0.78	0.78	73
13	0.85	0.82	0.84	97
14	0.75	0.73	0.74	78
15	0.72	0.49	0.58	96
16	0.79	0.69	0.74	71
17	0.71	0.71	0.71	59
18	0.97	0.70	0.81	80
19	0.90	0.77	0.83	71
20	0.52	0.89	0.66	76
21	0.58	0.77	0.66	69
22	0.81	0.85	0.83	66
23	0.76	0.82	0.79	71
24	0.80	0.65	0.72	69
25	0.90	0.68	0.77	77
26	0.48	0.56	0.51	79
27	0.96	0.79	0.87	86
28	0.91	0.81	0.85	83
29	0.73	0.58	0.64	83
accuracy			0.72	2250
macro avg	0.75	0.72	0.72	2250
weighted avg	0.75	0.72	0.73	2250

	precision	recall	f1-score	support
0	0.96	0.84	0.90	62
1	0.54	0.44	0.48	73
2	0.74	0.89	0.81	63
3	0.54	0.79	0.64	63
4	0.65	0.47	0.55	74
5	0.81	0.87	0.84	63
6	0.92	0.91	0.91	65
7	0.95	0.90	0.92	78
8	0.90	0.98	0.93	81
9	0.88	0.94	0.91	83
10	0.79	0.90	0.84	78
11	0.95	0.90	0.93	83
12	0.92	0.90	0.91	73
13	0.92	0.95	0.93	97
14	0.82	0.83	0.83	78
15	0.81	0.81	0.81	96
16	0.87	0.82	0.84	71
17	0.85	0.85	0.85	59
18	0.96	0.88	0.92	80
19	0.90	0.86	0.88	71
20	0.95	0.79	0.86	76
21	0.89	0.78	0.83	69
22	0.95	0.88	0.91	66
23	0.87	0.93	0.90	71
24	0.76	0.78	0.77	69
25	0.85	0.90	0.87	77
26	0.55	0.58	0.56	79
27	0.91	0.92	0.91	86
28	0.89	0.93	0.91	83
29	0.86	0.81	0.83	83
accuracy			0.84	2250
macro avg	0.84	0.83	0.83	2250
weighted avg	0.84	0.84	0.84	2250

CONCLUSION



The enhanced ResNet model demonstrated substantial improvement over the baseline, achieving an 84% accuracy on the test dataset. The deployment interface further showcases its practical utility.